

K-NEAREST NEIGHBOURS LABORATORY

Gabriele Francesco Berruti

University of Genoa
Via Roma 37/6
Genova, Italy
gabriele.berruti@unige.edu

Ibrahim Altufayli

University of Genoa
Via Roma 37/6
Genova, Italy
ibrahim.altufayli@unige.edu

ABSTRACT

This report centers on the k-NN algorithm due to its intuitive nature and practical effectiveness. The lab objective is to construct a KNN classifier and train it on synthetic 2D data for binary classification, incorporating noise to assess its impact on performance. The lab entails visualizing how the classifier separates data in 2D space and computing misclassification rates. The evaluation section focuses on comparing noisy and clean datasets, exploring varying noise levels and adjusting the number of neighbors (k). The experiments showed something important: When there's a lot of noise in the data, we need to use a higher 'k' value to estimate a more complicated decision boundary. But when the data is cleaner, we can use a smaller 'k' value to estimate a simpler function. This means the k-NN method can adjust to different types of data, which is really useful in practical situations where data might not always be perfect.

1. INTRODUCTION

Classification is a fundamental task in machine learning that involves sorting data into categories. It's like teaching a computer to recognize patterns and decide which group new data belongs to. This process is essential in various applications, such as spam detection, image recognition, and medical diagnosis. Several methods have been developed for classification, each with its own strengths and weaknesses. Simple methods like Decision Trees, Naive Bayes, and Logistic Regression are widely used due to their ease of implementation and interpretability. However, among these, the k-Nearest Neighbors (k-NN) algorithm stands out for its simplicity and effectiveness. k-NN is a type of instance-based learning, where the algorithm makes decisions based on the closest data points in the training set. When a new data point needs to be classified, k-NN looks at the 'k' nearest neighbors and assigns the most common category among them to the new point. This approach makes k-NN intuitive and easy to understand. In this report, we focus on k-NN for its straightforward mechanism and practical effectiveness. The goal of this lab is to build a KNN classifier and train

it on a synthetic 2D dataset for binary classification that we create. We will then add noise to the data to see how it impacts the KNN's performance, making the task more challenging. The lab includes visualizing how the classifier separates the data in 2D space and calculating the percentage of incorrectly classified points out of the total samples. After setting everything up, we will run experiments to compare how well the classifier works with both clean and noisy data, by changing the amount of noise and the value of 'k'.

2. ALGORITHMS AND METHODS

In this section we will describe the different Algorithms and Methods have been implement or used in the lab.

2.1. Dataset generation for binary classification

We start by generating a training set for binary classification problems. This involves creating random 2D points on a plane and assigning them binary labels based on their position relative to a linear separator. The function `linearBinaryClass` takes a sample size `n`, lower and upper bounds "low D", "high D" for the domain of the samples, and the linear function parameters `m`, `q` to generate a binary classification dataset, returning `X`(consists of 2-dimensional samples) and `Y` (consists of 1-dimensional binary labels).

2.2. Computing the distance between input points

To build the KNN estimator, we need to use a distance function. We implemented a function that calculates the Euclidean distance between two points in 2D space.

2.3. Adding noise to the samples

To make the task harder, we may want to perturb the labels with some noise. In our case, we have binary labels and a common way of adding noise is to flip the value of a small percentage of the labels. For example, if a label was `a+b` it will become `a-b`.

2.4. The KNN classifier

It takes four parameters: X_{tr} (training data inputs), Y_{tr} (training data labels), k (number of nearest neighbors to consider), and X_{te} (test data inputs). First, it checks if the labels in Y_{tr} are either +1 or -1, raising an exception if not. It then adjusts k if it exceeds the number of training samples. The function initializes an array Y_{pred} to store the predicted labels for the test data points. Next, it calculates the Euclidean distances between each test sample in X_{te} and all training samples in X_{tr} . For each test sample, it retrieves the distances to all training samples, sorts them, and selects the indices of the k nearest neighbors. It then computes the average label of these nearest neighbors and assigns the sign of this average as the predicted label for the test sample. Finally, the function returns the array Y_{pred} containing the predicted labels for the test data points.

2.5. Visualizing the separating function

visualizes the separating function on the training set, which is the function estimated by the classification algorithm to discriminate between classes. It takes three parameters: X_{tr} (the matrix of training set inputs), Y_{tr} (the array of training set labels), and k (the number of neighbors to consider). It interpolates the predicted labels (Y_{pred}) to create a continuous representation of the separating function, then plots the contour lines indicating the decision boundaries.

2.6. Evaluating the goodness of a classifier

To evaluate how good is the classification function estimated by the KNN, we compare the predicted binary labels and expected (true) ones. It then counts the number of elements where the predicted label is not equal to the true label, to calculate the error rate. This error rate represents the proportion of misclassified samples out of all samples in the dataset.

3. EXPERIMENTAL ANALYSIS

3.1. Experiment with not-noisy data

In the first section we will first evaluate the behavior of the KNN classifier on the training set and the test set in the case in which the creation of the two datasets will take place without noise. In particular, in this phase we specify the parameters relating to the number of samples to be generated, the lower and upper bounds for the domain of the samples and the parameters relating to the linear function (1) which will allow the data to be generated. At the end of this procedure we will obtain 2-dimensional samples (X) associated with 1-dimensional binary labels (Y), where the Y labels represents the real values linked to the training set. One

of the metrics that we will adopt in this section involves the calculation of the error between the predictions made by the model trained on the training set and the test set (or the training set itself). The first experiments we carry out in this phase is to evaluate the classifier on the training set with a parameter k equal to 1. So after training the model on the training set we evaluate it on the same set. As you can imagine, the error obtained is equal to zero since the values in the two set are literally the same. In the second experiment, similarly to what we did previously, we carry out a comparison between two sets where however they present different values (even if the generation is still described by a linear function). The results obtained by using and comparing the training set and the test set with $k=1$ allow us to obtain an error of 0.03, which means that 3 percent of the predictions made on the test set do not correspond to the real values. This behavior is due to the fact that using such a low parameter k does not allow us to extrapolate the form of the function since we restrict the search to a single (nearby) element and not to a series as we will see later.

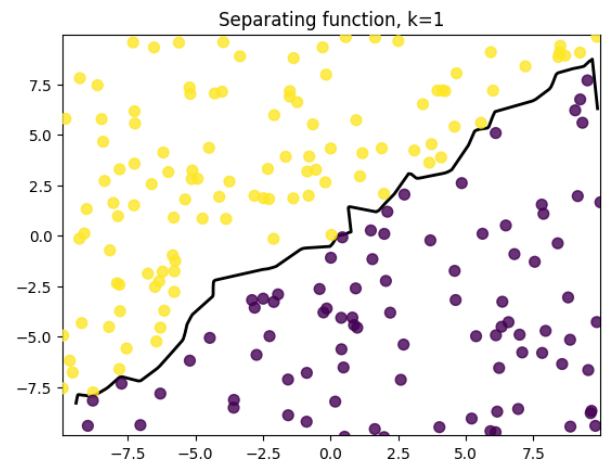


Figure 1: Example of a non-noisy dataset

3.2. Experiment with Noisy data and different K

This section will deal (as can be understood from the title) with the comparison between two sets (training and test set similarly to before) in which, however, the generation of data will be disturbed by random noise which does not allow the data to be distributed according to only a linear function with different parameters. Specifically, randomness mainly refers to which of the two labels to associate the pair of values. The quantity of pairs (in percentage) at which the labels will be modified is determined by the parameter P while the quantity of neighbors to take into consideration is always managed by k . In the first experiment of this section, a value of 10 is associated with the parameter P (which in-

icates changing 10 percent of the labels) while k remains equal to 1 (from the previous experiment). From this experiment we can see how the error obtained on the training set continues to remain equal to zero since the algorithm will once again compare two identical sets. The situation however varies when we feed the test set to the algorithm, in this case the calculated error is 0.19, which indicates that approximately 20 percent of the predicted labels do not reflect the real labels. Making a comparison with the experiment without noise, we note that the error increases considerably. By fixing the parameters decided previously and modifying exclusively the value of k we can notice that the results vary considerably. For example, using a K equal to 5, and therefore evaluating the classifier on the training set itself, we notice an unusual behavior, i.e. the error on the predictions is higher than that where k was equal to 1. Varying from 0 to 0.115. This is because observing a larger spectrum of nearby pairs can lead to a prediction that is different from the actual label. Furthermore, it is useful to note that the percentage error value is comparable to the value of labels that are "flipped", this is because the "noisy" pairs will most likely be "immersed" in a neighborhood composed of pairs belonging to the other label. The situation is reversed when we evaluate the noisy data on the test set. In this case the error decreases compared to the previous case, varying from 0.19 to 0.145. This is due to the fact that, even if the data are noisy, the values that are along or near the separation function will have a better spectrum of elements to compare with will be able to determine their label in a more coherent way while for the elements in the test sets (as in the training set) that are surrounded only by labels that differ from the real one there is no possibility that they will be associated correctly. Finally, the last experiment involves the use of a k equal to 10. This value allows us to increase the number of neighboring pairs with which it is possible to compare. In this specific situation the choice of the aforementioned parameter allows the classifier to obtain a slightly worse result than the error relating to the k equal 5, since the error between the predicted value and the real one varies from 0.115 to 0.135 while the calculation of the error on the test part goes from 0.145 to 0.11. These results depend a lot on the random generation of the data but it can still be noted that a smaller k causes the classifier to fall in love with the data (overfitting the training set) while increasing k in general tends to generalize more because it searches in a wider area.

4. CONCLUSIONS

The experiments clearly demonstrate that the choice of ' k ' in the k -NN algorithm depends on the level of noise in the data. When dealing with noisy data, a higher ' k ' is neces-

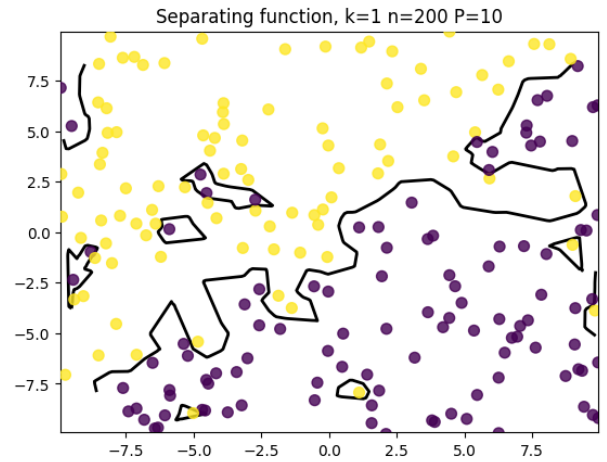


Figure 2: Example of a noisy dataset

sary to build a more complex decision boundary that can effectively separate signal from noise. Conversely, in cleaner datasets, lower ' k ' values are enough to capture the underlying patterns in the data without being overly influenced by noise. This showed the importance of tailoring the choice of ' k ' to the specific characteristics of the dataset, optimizing the algorithm's performance for different levels of data quality.

5. EQUATIONS

Equations of the linear function : (1) [1],

$$f(x) = m(x) + q \quad (1)$$

where $f(x)$ is the function that allows us to sub-divide data (by associating a label) based on the parameters m and q which represents the angular coefficient and the coefficient q is a known term .

6. REFERENCES

- [1] E. Williams, *Fourier Acoustics: Sound Radiation and Nearfield Acoustic Holography*. London, UK: Academic Press, 1999.